



TALK-TO-DATA (T2D)

Phase 1 Annexes: Framing

Daniel Brulé

[LinkedIn](#) - [GitHub](#)

Version 1.0 - June 2026

Table of Contents

1.	How to use this annex pack	3
2.	Activity question banks	3
2.1.	Business framing	3
2.2.	User and workflow discovery	3
2.3.	Initial data landscape scan	4
2.4.	Initial semantic framing	5
2.5.	Security and governance framing	5
2.6.	Solution architecture framing.....	6
2.7.	Delivery planning and MVP boundary.....	7
2.8.	Evaluation design framing	7
2.9.	Operating model framing	8
3.	Activity output templates	9
3.1.	Business framing	9
3.2.	User and workflow discovery	10
3.3.	Initial data landscape scan	11
3.4.	Initial semantic framing	11
3.5.	Security and governance framing	12
3.6.	Solution architecture framing.....	12
3.7.	Delivery planning and MVP boundary.....	13
3.8.	Evaluation design framing	14
3.9.	Operating model framing	14
4.	Activity scorecards.....	15
4.1.	Business framing scorecard	15
4.2.	User and workflow discovery scorecard.....	15
4.3.	Initial data landscape scan scorecard	16
4.4.	Initial semantic framing scorecard	16
4.5.	Security and governance framing scorecard	16
4.6.	Solution architecture framing scorecard	17
4.7.	Delivery planning and MVP boundary.....	17
4.8.	Evaluation design framing scorecard	17
4.9.	Operating model framing scorecard	18
5.	Overall phase readiness scorecard	19
6.	Risk, decision, assumption and dependency logs.....	20

6.1.	Risk log	20
6.2.	Decision log	21
6.3.	Assumption log	22
6.4.	Dependency log.....	23

1. How to use this annex pack

This annex pack contains detailed tools and templates supporting the Phase 1 Framing Guide. It is intended for facilitation, documentation and project delivery. The main guide explains the framing logic; this annex provides the reusable working material.

2. Activity question banks

2.1. Business framing

Area	Questions
Strategic objective	Why Talk-to-Data now? Is the driver productivity, self-service, faster reporting, executive insight, analyst augmentation or another business need?
Priority domain	Which domain should be addressed first: finance, sales, operations, customer, risk, product or another area?
Scope	What should the assistant answer in v1? What is explicitly out of scope? What is the longer-term ambition?
Users	Who is the first user group: executives, managers, analysts, frontline teams or another group?
Timeline	What is the expected timeline for discovery, POC, MVP, pilot and production?
Risk appetite	What level of risk is acceptable for answer accuracy, data exposure, user autonomy and business impact?
Sponsorship	Who funds delivery? Who owns the business outcome?
Governance	Who approves scope, metric definitions, access rules, risk posture, pilot and go-live? What is reported, to whom and when?
Budget	What budget envelope is available for discovery, POC/MVP, production readiness and ongoing run costs?
Success	What decisions will be improved? What time, cost, quality or risk benefit is expected?
Failure impact	What would a harmful or unacceptable answer look like?

2.2. User and workflow discovery

Area	Questions
User groups	Who are the target users: executives, managers, analysts, frontline teams or external users?
User volume and location	How many users are expected? Where are they located? Are they office-based, remote, mobile or global?
Skill level	Are users data-literate? Can they interpret metrics, SQL, charts, confidence levels and caveats?
Access context	Which devices and channels do they use: desktop, mobile, Teams, Slack, BI tool, intranet or another interface?
Current workflow	How do users answer these questions today: dashboard, analyst request, SQL, Excel, BI tool or manual process?

Workflow frequency	How often are questions asked: daily, weekly, monthly, ad hoc or board-cycle driven?
Pain points	Is the issue speed, trust, lack of data skills, fragmented tools, inconsistent metrics, analyst dependency or slow turnaround?
Decision context	What decisions are made from the answers? Are they operational, tactical, strategic, regulatory or investment-related?
Question types	Are users asking KPI, trend, comparison, variance, root-cause, forecasting, explanation or data-quality questions?
Question complexity	Are questions simple lookups, aggregations, multi-step analysis, cross-domain joins or ambiguous business questions?
Follow-ups	Do users need conversational drill-downs, filters, comparisons, rephrasing or “why has this changed?” analysis?
Output format	Do users expect a number, chart, table, SQL, narrative, export, dashboard link or recommended action?
Trust requirements	Do users need source tables, metric definitions, SQL, confidence level, caveats or lineage to trust the answer?
Human review	When should the assistant answer directly, ask clarification, escalate to an analyst or refuse?
Error tolerance	What level of error is acceptable by question type? Which answers require validation before being shown?

2.3. Initial data landscape scan

Area	Questions
Sources	Which systems contain the required data: warehouse, lakehouse, BI semantic model, operational systems, APIs or files?
Data ownership	Who owns each source, table, metric and business definition?
Trusted assets	Which tables, views, dashboards or marts are certified, trusted, deprecated or experimental?
Current usage	Which tables and dashboards are currently used to answer the priority questions?
Data grain	At what level are facts stored: customer, transaction, account, day, product, contract or event?
Freshness	How frequently is the data updated: real-time, daily, weekly, monthly or manually refreshed?
Coverage	Does the available data cover the required users, geographies, products, time periods and entities?
Lineage	Can key metrics be traced from report back to model, table and source system?
Data quality	Are there known issues with completeness, duplicates, nulls, timeliness, reconciliation or reliability?
Existing semantic assets	Are there existing BI models, certified datasets, dbt models, cubes or semantic layers that can be reused rather than rebuilt?
Data change risk	Are source systems, tables, metrics or ownership expected to change during the delivery timeline?

2.4. Initial semantic framing

Area	Questions
Metrics	Are key metrics clearly defined: revenue, margin, churn, active customer, pipeline, conversion?
Consistency	Are definitions consistent across teams, geographies and reporting packs?
Dimensions	Which dimensions are approved for slicing: time, region, customer, product, channel, segment, business unit?
Joins	Which joins are likely to be approved and safe? Are there ambiguous or many-to-many joins to avoid?
Filters	What standard exclusions apply: test accounts, cancelled orders, internal users, inactive products, intercompany transactions?
Business glossary	Are business terms, synonyms and abbreviations documented?
Calculation logic	Are formulas documented in SQL, BI, dbt, semantic layer or only in analysts' heads?
Query examples	Are trusted SQL examples or BI calculations available for grounding and evaluation?
Ambiguity	Which terms require clarification before answering?
Certified answer sources	Which dashboards, reports, SQL queries or finance packs are considered official?

2.5. Security and governance framing

Area	Questions
Identity	How is the user authenticated: SSO, BI permissions, warehouse roles or app-level identity?
Authorisation	Should T2D inherit existing user permissions, use service-account permissions or apply a separate policy layer?
Row-level security	Can users only see their region, business unit, portfolio, customer segment or assigned accounts?
Column-level security	Are salary, PII, margin, customer identifiers, contract terms or sensitive fields masked or blocked?
Aggregation risk	Could users infer restricted data from aggregates, small cohorts, drill-downs or repeated queries?
Data classification	Which datasets contain public, internal, confidential, restricted, PII, financial or regulated data?
Compliance	What constraints apply: GDPR, financial controls, client confidentiality, regulatory requirements, data residency or retention?
Model / vendor policy	Can data or metadata be sent to an external LLM provider? Are private endpoints, zero-retention or on-prem models required?
Logging	What must be logged: user, question, context, SQL, result summary, answer, errors and feedback?
Audit	Who can review past questions, SQL, answers and access events? At what cadence and for what purpose?
Retention	How long should prompts, SQL, results and feedback be retained? What must be redacted or excluded?

2.6. Solution architecture framing

Area	Questions
Cloud and platform landscape	Which cloud providers, data platforms, BI tools, identity providers and integration platforms are currently in use?
Technology standards	Are there mandated tools, preferred vendors, approved LLM providers, architecture principles or prohibited technologies?
Architecture governance	What is the architecture review process? Who approves solution design, security design and production deployment?
Architecture approval	What evidence is required for architecture approval: high-level design, data flow, security model, integration diagram, non-functional requirements, cost estimate or support model?
IT roadmap	Are there planned changes: cloud migration, data platform migration, BI migration, identity changes or governance tooling?
Interface	Should T2D live in a dedicated UI, BI tool, Teams/Slack, internal portal, notebook, CRM or workflow application?
User experience	Should the interface support chat history, follow-up questions, charts, SQL visibility, citations, exports, feedback or escalation?
Model choice	Which LLM provider or model is allowed? If none, is there a process to select one?
Model routing	Should different models be used for intent classification, metadata retrieval, SQL generation, validation and explanation?
Context strategy	What instructions, examples, metadata, business rules and user context should be retrieved?
Data connectivity	Which query engine, warehouse, semantic layer or API will be used?
SQL validation	Should validation use SQL parser, dry-run, explain plan, policy engine, allow-listed tables, row limits or second model review?
Safety controls	What should be blocked: SELECT *, PII columns, raw exports, high-cost queries, small cohorts or destructive commands?
Deployment model	Should deployment use cloud, VPC, private endpoint, containers, Kubernetes, serverless, managed AI platform or internal platform?
Observability	Where are latency, cost, token usage, errors, SQL, feedback, quality and security events logged?
Non-functional requirements	What latency, scalability, reliability, cost, auditability and maintainability requirements apply?

2.7. Delivery planning and MVP boundary

Area	Questions
MVP boundary	Which users, domains, questions, datasets and answer types are included in MVP?
Exclusions	What is explicitly excluded from MVP?
Backlog	What are the initial epics and work packages?
Timeline	What is the timeline for discovery, readiness, governed layer, architecture, prototype, security, evaluation, pilot and production?
Team capacity	Which business, data, engineering, security, SME and product roles are available?
Budget	What budget is available for delivery effort, data engineering, LLM usage, warehouse compute, tooling, security review and support?
Approval gates	What decisions are required before discovery, prototype, pilot and production?
Delivery risks	What could delay or block MVP delivery?
Dependencies	Which data, access, architecture, evaluation or governance dependencies are critical?
Delivery method	Will delivery follow agile, phased gates, discovery sprint, POC, pilot or another approach?

2.8. Evaluation design framing

Area	Questions
Evaluation scope	Are we evaluating SQL correctness, business correctness, safety, answer quality, latency, cost, adoption or all of these?
Success bar	What level of accuracy, safety, latency and user satisfaction is acceptable before MVP, pilot and production?
Golden questions	What are the 30–100 questions users actually expect the assistant to answer?
Ground truth ownership	Who can confirm expected SQL, metric logic, final answer and caveats?
Evaluation dimensions	What counts as correct: valid SQL, right metric, right grain, right filters, correct number, useful explanation?
Safety evaluation	What data must never be exposed? What access-control, PII, aggregation and inference risks must be tested?
Operational metrics	What latency, cost per question, error rate and timeout rate are acceptable?
User feedback	Can users mark answers as correct, wrong, unclear, unsafe or not useful?
Regression approach	Will the golden set be rerun after prompt, model, semantic layer, schema or code changes?
Evaluation ownership	Who owns evaluation after launch: product, analytics, QA, SMEs, data governance or platform team?
Failure taxonomy	Are failures due to intent, metadata retrieval, SQL generation, semantic mismatch, data quality, permissions, UX or model explanation?

2.9. Operating model framing

Area	Questions
Product ownership	Who owns the T2D product roadmap after launch?
Business ownership	Who owns the business outcome and prioritises new use cases?
Semantic ownership	Who approves changes to metrics, dimensions, joins, filters and business definitions?
Data ownership	Who owns source data quality, freshness and lineage?
Technical ownership	Who maintains the application, orchestration, prompts, connectors and deployment?
Security ownership	Who maintains access rules, masking, row-level security and audit controls?
Evaluation ownership	Who maintains golden questions, regression tests and quality thresholds?
Support ownership	Who triages user issues: data team, BI team, platform team, product owner or support desk?
Funding	Who funds discovery, build, production deployment and ongoing run costs?
Change management	What communication, training and adoption support is needed?
Workforce impact	How will T2D affect analysts, reporting teams and business users?
Continuous improvement	How will new questions, failed answers, semantic gaps and user feedback be prioritised?

3. Activity output templates

These columns are illustrative examples only. They are intended to show the expected structure and level of detail.

3.1. Business framing

Item	Description	Owner ¹	Status ¹	Notes ¹
Problem statement	Why Talk-to-Data is being considered	Product owner / sponsor	Draft / agreed	Should link to a specific business pain point
Priority domain	First business area in scope	Business lead	Draft / agreed	Should avoid multi-domain MVP unless justified
MVP scope	First users, questions, datasets and excluded areas	Product owner	Draft / agreed	Include explicit exclusions
Strategic value case	Expected benefit and decision improvement	Sponsor	Draft / agreed	Should be measurable where possible
Risk appetite	Acceptable risk posture for MVP	Sponsor / risk owner	Draft / agreed	Must inform security and evaluation
Governance route	Decision forums and approval gates	Delivery lead	Draft / agreed	Include pilot and go-live approval
Budget envelope	Funding for discovery, MVP and run costs	Sponsor / finance	Draft / agreed	Include data, AI, platform and support costs

3.2. User and workflow discovery

User group	Workflow / decision	Current process	Priority questions	Frequency	Pain point	Expected answer format	Trust requirement	MVP relevance
Executive team	Performance review	Board pack and follow-up analyst requests	Revenue, margin, target variance, regional performance	Monthly / ad hoc	Slow follow-up questions	KPI, chart, narrative	Source, metric definition, caveat	High
Sales managers	Pipeline and account review	CRM dashboard and analyst extracts	At-risk opportunities, declining customers, conversion trends	Weekly	Fragmented reporting	Table, chart, explanation	Source, filters, freshness	High
Analysts	Exploratory analysis	SQL, BI, Excel	Root-cause, variance, drill-downs	Daily / weekly	Manual query effort	SQL, table, export	SQL visibility, lineage	Medium
Operations managers	Service or backlog review	Dashboard and spreadsheets	Backlog, delays, location variance	Daily / weekly	Limited self-service	KPI, table, chart	Freshness, definitions	Medium

3.3. Initial data landscape scan

Priority question area	Candidate source / report	System	Owner	Trusted status	Coverage	Freshness	Known issue	Step 2 validation needed
Revenue performance	Finance revenue mart / finance pack	Warehouse / BI	Finance Analytics	Likely certified	Region, product, month	Daily / monthly close	Definition may vary by report	Confirm source of truth and metric logic
Sales pipeline	Sales pipeline dashboard	CRM / BI	Sales Ops	Trusted reference	Opportunity, owner, stage	Daily	Stage definitions may vary	Confirm grain, filters and owner
Customer spend	Revenue mart + customer master	Warehouse	Data / Sales Ops	Partially trusted	Customer, month	Daily	Customer duplicates	Validate joins and customer hierarchy
Churn	CRM + billing	CRM / Billing	Sales Ops / Finance	Unknown	Customer, segment, month	Unknown	No agreed definition	Confirm definition and source ownership

3.4. Initial semantic framing

Metric / term	Possible meaning	Default framing assumption	Source of truth candidate	Owner	Ambiguity risk	Step 2 validation needed
Revenue	Gross, net, recognised, booked	Net revenue for MVP unless sponsor decides otherwise	Finance mart / finance pack	Finance	High	Agree approved definition and filters
Customer	CRM account, billing customer, legal entity	Billing customer for financial questions	Customer master	Sales Ops / Data	High	Confirm hierarchy and joins
Churn	No spend, cancelled contract, lost account	Not agreed	CRM + billing	Sales Ops	High	Define churn and ownership
Pipeline	Open opportunities, weighted pipeline, forecast	Sales pipeline by approved CRM stage	Sales BI dashboard	Sales Ops	Medium	Confirm stage logic and exclusions

3.5. Security and governance framing

Area	Initial assumption	Risk level	Owner	Validation needed	Decision required
Authentication	Use enterprise SSO	Medium	Security architect	Confirm integration route	Yes
Authorisation	Inherit existing BI / warehouse permissions where possible	High	Security / data platform	Confirm feasibility and gaps	Yes
Sensitive data	PII excluded from MVP unless specifically approved	High	DPO / governance	Confirm data classification	Yes
Row-level access	Apply region / business-unit restrictions for business users	High	Data owner / security	Confirm available access model	Yes
SQL visibility	Analysts only for MVP	Medium	Product owner / security	Confirm user groups	Yes
Logging	Log user, question, retrieved context, generated SQL, result status and feedback	Medium	Platform owner	Confirm retention and redaction	Yes

3.6. Solution architecture framing

Area	Initial architecture assumption	Constraint / dependency	Owner	Validation needed	MVP implication
User interface	Start with controlled internal UI or Teams/Slack integration	Enterprise UX and security approval	Product / platform	Confirm preferred channel	Defines user access path
Model provider	Use approved enterprise LLM provider	Vendor and data policy	AI architect / security	Confirm permitted model and endpoint	May affect cost and latency
Data access	Query governed marts or semantic layer only	Step 2 data readiness	Data architect	Confirm available query layer	Limits MVP scope
SQL validation	Use parser, allow-listed schemas, dry-run and row limits	Platform tooling	AI architect	Confirm validation approach	Required before pilot
Observability	Log prompts, SQL, result status, errors, feedback and cost	Logging tools and policy	Platform owner	Confirm retention and redaction	Required for safe MVP
Deployment	Use approved enterprise deployment pattern	Architecture approval	Platform architect	Confirm environment path	Defines timeline

3.7. Delivery planning and MVP boundary

Workstream	Key activity	Owner	Dependency	Indicative timing	MVP criticality	Notes
Product	Confirm MVP questions and users	Product owner	Sponsor decision	Week 0–1	High	Required before Step 2
Data	Validate candidate data assets	Data architect	Access to sources	Week 1	High	Feeds Step 2
Semantic	Draft metric and business definitions	Analytics engineer	SME availability	Week 1	High	Feeds governed layer
Security	Confirm access and exposure constraints	Security architect	Data classification	Week 1–2	High	Required before prototype
Architecture	Confirm target pattern	AI architect	Platform standards	Week 1–2	High	Required before build
Evaluation	Create golden question plan	AI architect / QA	Question backlog	Week 1–2	Medium	Required before pilot
Delivery	Build MVP roadmap and backlog	Delivery lead	All workstreams	Week 1	High	Phase-gate input

3.8. Evaluation design framing

Evaluation item	Definition	Owner	MVP expectation	Evidence needed
Golden question set	Representative business questions for MVP	Product owner / business analyst	30–50 initial questions	User interviews and existing reporting requests
Ground truth	Expected answer, SQL, metric logic and caveats	SME / analytics engineer	Available for priority questions	Existing reports, SQL and SME sign-off
Correctness criteria	SQL, metric, grain, filter and answer accuracy	QA / AI architect	Defined before pilot	Evaluation rubric
Safety criteria	Access, PII, small cohort and restricted field tests	Security / governance	Defined before prototype	Security test cases
Operational metrics	Latency, cost, error rate and timeout thresholds	Platform owner	Indicative targets	Monitoring plan
Feedback model	User feedback categories and triage route	Product owner	Defined for pilot	Feedback workflow

3.9. Operating model framing

Ownership area	Proposed owner	Responsibility	MVP need	Open question
Product	Product owner	Roadmap, scope, prioritisation and user experience	Required	Confirm long-term owner
Business	Business sponsor	Business value and decision outcomes	Required	Confirm sponsor and forum
Semantic	Metric owners / analytics	Metrics, definitions, joins, filters and glossary	Required	Confirm decision rights
Data	Data owners / data engineering	Quality, freshness, lineage and source changes	Required	Confirm operational SLA
Technical	Platform / AI owner	App, prompts, orchestration, connectors and deployment	Required	Confirm support model
Security	Security / governance	Access, masking, audit and compliance controls	Required	Confirm approval path
Evaluation	QA / analytics / SMEs	Golden questions, tests and thresholds	Required	Confirm ongoing owner
Support	Product / BI / platform support	Issue triage and user support	Required for pilot	Confirm escalation route

4. Activity scorecards

4.1. Business framing scorecard

Criterion	Green	Amber	Red
Business problem	Clear, specific and valuable	Broad but directionally clear	Generic or technology-led
MVP domain	Single priority domain agreed	Domain likely but not confirmed	Multiple competing domains
Scope boundary	In-scope and out-of-scope clear	Some boundary questions remain	Scope open-ended
Sponsor ownership	Sponsor and business owner confirmed	Sponsor identified but role unclear	No accountable sponsor
Success criteria	Measurable outcomes defined	Outcomes qualitative only	No success criteria
Budget envelope	Budget range and approval route clear	Budget assumption only	No funding view

4.2. User and workflow discovery scorecard

Criterion	Green	Amber	Red
Target users	First user group clearly defined	Multiple candidate groups remain	User group unclear
Workflow fit	Clear workflow and decision context	Workflow partly understood	No clear workflow
Question backlog	30–50 candidate questions captured	Initial questions captured	Few or vague questions
Frequency and value	Frequent or high-value questions	Value plausible but unquantified	Low frequency or unclear value
Output expectations	Answer format and trust needs clear	Some uncertainty remains	Expectations undefined
Human review needs	Escalation and refusal cases identified	Some review needs unclear	No control model

4.3. Initial data landscape scan scorecard

Criterion	Green	Amber	Red
Candidate sources	Likely sources identified for most priority questions	Sources identified for some questions	Sources unclear
Ownership	Business and technical owners known	Some owners missing	Owners unclear
Trusted assets	Certified or trusted assets likely available	Mixed trusted and informal assets	No trusted assets identified
Coverage	Coverage appears sufficient for MVP	Some coverage questions remain	Material coverage gaps likely
Freshness	Freshness likely adequate	Freshness uncertain	Freshness likely unsuitable
Known data risks	Key risks captured for Step 2	Some risks captured	Risks unknown

4.4. Initial semantic framing scorecard

Criterion	Green	Amber	Red
Key metrics	Priority metrics identified and likely owned	Metrics identified but ownership unclear	Metrics unclear
Business terms	Main ambiguous terms captured	Some ambiguity remains	Ambiguity not assessed
Dimensions	Key dimensions identified	Some dimensions uncertain	Dimensions unclear
Filters	Likely standard exclusions known	Filters partly known	No filter assumptions
Certified sources	Candidate official sources identified	Multiple sources compete	No source of truth
Semantic risk	Material risks visible for Step 2	Some risks captured	Semantic risk unknown

4.5. Security and governance framing scorecard

Criterion	Green	Amber	Red
User access model	Existing access model likely reusable	Adaptation needed	No access model
Sensitive data	Sensitive data known and manageable	Classification incomplete	Sensitivity unknown
Compliance constraints	Key constraints understood	Some constraints pending	Compliance not engaged
Audit needs	Logging and audit needs clear	Partial audit view	Audit unclear
Exposure controls	High-level controls agreed	Some controls unresolved	No safe exposure model
Approval route	Security and risk owners engaged	Owners identified but not engaged	No approval path

4.6. Solution architecture framing scorecard

Criterion	Green	Amber	Red
Platform fit	Existing platforms can support MVP	Some platform gaps	No viable platform route
Model policy	Approved LLM route clear	Approval pending	No approved model route
Data connectivity	Query route likely available	Connectivity uncertain	No safe query path
Validation controls	SQL and safety controls feasible	Controls need design	Controls unclear
Deployment route	Environments and release route known	Some governance pending	No deployment path
Non-functional view	Latency, cost, reliability and audit needs captured	Partial view	Not assessed

4.7. Delivery planning and MVP boundary

Criterion	Green	Amber	Red
MVP boundary	Clear and realistic	Some scope uncertainty	Scope too broad or unclear
Timeline	Plausible timeline agreed	Timeline tight or assumption-led	Timeline unrealistic
Team capacity	Required roles available	Some roles constrained	Critical roles unavailable
Budget	Budget envelope and approval route clear	Budget indicative only	No budget route
Backlog	Initial backlog structured	Backlog incomplete	No delivery backlog
Governance gates	Decision points clear	Some gates unclear	No approval path

4.8. Evaluation design framing scorecard

Criterion	Green	Amber	Red
Golden questions	Representative set defined	Initial set incomplete	No question set
Ground truth	Owners and validation route clear	Some ownership gaps	No validation owner
Quality dimensions	Scoring criteria defined	Criteria partial	No evaluation approach
Safety tests	Safety scenarios identified	Some scenarios missing	Safety not assessed
Regression approach	Rerun approach defined	Manual approach only	No regression plan
Feedback loop	Feedback capture and triage clear	Partial feedback route	No feedback process

4.9. Operating model framing scorecard

Criterion	Green	Amber	Red
Product ownership	Clear owner and roadmap accountability	Interim owner only	No owner
Business ownership	Sponsor and outcome owner clear	Sponsor unclear	No business owner
Semantic ownership	Metric decision rights clear	Some owners missing	No semantic ownership
Technical ownership	Platform and application owner clear	Split ownership unresolved	No support owner
Security ownership	Security and audit responsibilities clear	Some responsibilities unclear	No security owner
Support model	Triage and escalation model defined	Informal support only	No support route
Adoption approach	Training and communication approach defined	Adoption assumptions only	No adoption plan

5. Overall phase readiness scorecard

Readiness area	Rating	Evidence	Key risk	Action required	Owner
Business value readiness	Green / Amber / Red	Problem statement, value case, success criteria	Weak value case	Refine use case or pause	Sponsor / product owner
MVP scope readiness	Green / Amber / Red	Domain, users, included / excluded areas	Scope too broad	Narrow MVP boundary	Product owner
User and workflow readiness	Green / Amber / Red	User groups, workflows, priority questions	Low adoption or weak fit	Validate with users	Product owner / BA
Data landscape readiness	Green / Amber / Red	Candidate sources and ownership	Data not available or not trusted	Move to Step 2 validation	Data architect
Semantic readiness	Green / Amber / Red	Draft metrics, terms, filters, joins	Definitions inconsistent	Create semantic decision log	Analytics engineer
Security and governance readiness	Green / Amber / Red	Initial access, classification and audit assumptions	Exposure risk underestimated	Security review before prototype	Security architect
Architecture feasibility	Green / Amber / Red	Platform, model, integration and deployment assumptions	No approved technical route	Architecture review	AI architect
Evaluation readiness	Green / Amber / Red	Golden question plan and quality dimensions	No objective success measure	Define evaluation approach	AI architect / QA
Delivery readiness	Green / Amber / Red	Timeline, backlog, budget, capacity	MVP not deliverable	Re-plan or reduce scope	Delivery lead
Operating model readiness	Green / Amber / Red	Ownership, support and adoption assumptions	No sustainable run model	Define RACI and support route	Product owner
Overall framing readiness	Green / Amber / Red	Phase-gate decision pack	Proceeding with weak foundations	Proceed / pause / re-scope decision	Sponsor

Suggested rating logic:

Rating	Meaning
Green	Ready for governed data layer with limited caveats
Amber	Usable for MVP if specific gaps are resolved or accepted
Red	Not ready; material blocker exists
Grey	Not assessed yet

6. Risk, decision, assumption and dependency logs

6.1. Risk log

Risk	Impact	Likelihood	Mitigation	Owner	Status
T2D is pursued because of GenAI interest rather than a clear business need	Low adoption, weak value case, poor prioritisation	Medium	Anchor scope in priority decisions, workflows and measurable outcomes	Sponsor / product owner	Open
MVP scope is too broad	Delivery becomes unmanageable and answer quality drops	High	Start with one domain, one user group and bounded question set	Product owner	Open
Priority questions are vague or not frequent enough	T2D may not be the right solution	Medium	Validate questions with users and compare against dashboard/report alternatives	Business analyst	Open
No clear sponsor or decision owner	Delayed decisions and weak accountability	Medium	Confirm sponsor, decision rights and governance forum during framing	Delivery lead	Open
Data availability is assumed but not validated	Step 2 may reveal major blockers	Medium	Capture candidate sources and known gaps; validate in Step 2	Data architect	Open
Metrics are inconsistent across teams	Conflicting answers and low trust	High	Create semantic decision log and identify metric owners	Analytics engineer	Open
Security constraints are underestimated	Security approval may block MVP	Medium	Engage security, DPO and governance early	Security architect	Open
Budget excludes run and support costs	MVP may be unsustainable	Medium	Include platform, model, warehouse, support and semantic maintenance costs	Sponsor / finance	Open
Evaluation approach is left too late	Cannot prove correctness or safety	Medium	Define golden question plan during framing	AI architect / QA	Open
Operating model is unclear	Prototype cannot transition to live service	Medium	Define ownership and support assumptions before build	Product owner	Open

6.2. Decision log

Decision	Options	Recommended option	Decision owner	Due date	Status
Which domain should be used for MVP?	Finance / Sales / Operations / Customer / Product / Risk	Select one high-value, data-ready domain	Sponsor / product owner	End of framing	Open
Which user group should be included first?	Executives / managers / analysts / frontline users	Start with bounded group with clear workflow and support	Product owner	End of framing	Open
What is included in MVP?	Broad assistant / bounded question set / report replacement / analyst co-pilot	Bounded question set linked to specific decisions	Product owner	End of framing	Open
What is explicitly out of scope?	Cross-domain joins / PII / row-level exports / forecasting / agentic actions	Exclude high-risk and poorly governed areas from MVP	Product owner / security	End of framing	Open
What success measures apply?	Adoption / accuracy / time saved / cost per answer / satisfaction / trust	Use a balanced set of business, quality, safety and operational metrics	Sponsor / product owner	End of framing	Open
What is the approval path?	Sponsor only / steering group / architecture / security / data governance	Use proportionate phase gates for scope, security, pilot and go-live	Delivery lead	End of framing	Open
What budget is available?	Discovery only / MVP / pilot / production / run	Confirm budget across delivery and run	Sponsor / finance	End of framing	Open
Should T2D proceed to Step 2?	Proceed / pause / re-scope / stop	Proceed only if value, scope, ownership and feasibility are sufficient	Sponsor	End of framing	Open

6.3. Assumption log

Assumption	Impact if wrong	Validation method	Owner	Status
A high-value use case exists for T2D	Project may not justify investment	Confirm with sponsor and target users	Product owner	To validate
Priority users have recurring questions not well served by current BI	T2D may not be the right solution	User interviews and workflow review	Business analyst	To validate
Candidate data exists for priority questions	Step 2 may expose data blockers	Initial data landscape scan	Data architect	To validate
Existing reports or dashboards can provide starting logic	Metric assumptions may be wrong	Compare with BI owners and SMEs	Analytics engineer	To validate
Security constraints can be addressed within MVP	Prototype may be blocked	Security and DPO review	Security architect	To validate
Approved enterprise model and deployment route exists	Architecture may need rework	CTO/platform review	AI architect	To validate
SMEs are available to validate questions and answers	Evaluation may be weak	Confirm SME capacity	Product owner	To validate
Budget can cover delivery and run costs	MVP may not be sustainable	Finance review	Sponsor	To validate
Product and support ownership can be assigned	Capability may not scale after pilot	Operating model review	Delivery lead	To validate

6.4. Dependency log

Dependency	Needed by	Owner	Due date	Status	Impact if delayed
Sponsor and decision forum confirmed	Scope and proceed decision	Sponsor / delivery lead	End of framing	Open	Decisions delayed
Priority domain selected	User discovery and data scan	Product owner	End of framing	Open	Scope remains too broad
Target users identified	Workflow discovery and access review	Product owner	End of framing	Open	User needs unclear
Initial question backlog captured	Data and semantic readiness	Business analyst	End of framing	Open	Step 2 lacks focus
Candidate reports and dashboards identified	Data landscape scan	BI lead / data analyst	End of framing	Open	Source discovery delayed
Candidate data owners identified	Step 2 validation	Data architect	End of framing	Open	Ownership unclear
Security and governance stakeholders engaged	Access and exposure framing	Security architect	End of framing	Open	Late security blocker
Architecture standards confirmed	Solution feasibility	AI architect / CTO	End of framing	Open	Technical route unclear
Budget envelope confirmed	MVP planning	Sponsor / finance	End of framing	Open	Delivery cannot be committed
SME availability confirmed	Evaluation and semantic validation	Product owner	End of framing	Open	Ground truth and definitions delayed