



TALK TO DATA

Phase 5: Prototype / MVP build

Daniel Brulé

Brule.daniel@gmail.com

May 2026 - V1

Table of Contents

1. Executive summary	4
2. Phase overview	5
2.1. Objective.....	5
2.2. Scope of the phase	5
2.3. What this phase does not do	6
2.4. Expected duration and level of effort	6
2.5. Main participants and decision owners.....	6
3. Readiness decision and delivery implications	7
3.1. Possible Phase 5 outcomes.....	7
3.2. Minimum conditions to proceed.....	7
3.3. Common reasons to remediate, redesign or pause	8
3.4. How Phase 5 shapes later phases	8
4. MVP build activities overview.....	9
4.1. Activity sequence	9
4.2. Activity logic	10
4.3. Output need levels	10
4.4. Practitioner note.....	10
5. Core build activities	11
5.1. Confirm MVP scope and build plan.....	11
5.2. Confirm MVP governance, ownership and decision rights.....	12
5.3. Set up environment, CI/CD and test harness	13
5.4. Connect approved data, metadata and tools	14
5.5. Implement metadata retrieval and grounding.....	15
5.6. Implement AI/model routing and clarification	17
5.7. Implement query generation, validation and execution.....	18
5.8. Implement result interpretation and answer generation	19
5.9. Build lightweight UI and feedback capture	20
5.10. Tune, test and package evidence.....	21
6. Consolidated Phase 5 outputs.....	23
6.1. MVP evidence pack.....	23
6.2. Output quality test.....	23

7.	Exit criteria and handover	24
7.1.	Required exit outputs.....	24
7.2.	Handover to later phases	24
7.3.	Exit decision wording	24
7.4.	Practitioner note.....	25

1. Executive summary

Phase 5 turns the Phase 4 orchestration design into a bounded working MVP. If Phase 4 defines the intended question-to-answer flow, Phase 5 tests that flow against reality.

The purpose of this phase is not to build the final production service. It is to challenge the approved design through implementation: can the system retrieve the right metadata, select the right AI and model patterns, generate constrained queries, validate them before execution, interpret results, surface caveats, handle follow-up questions, fail safely and capture enough evidence for later validation?

The central discipline is to build for evidence, not demonstration value. A polished interface is not enough. The MVP should show which parts of the design work, which assumptions fail, where AI adds value, where deterministic controls are safer than model judgement, and where the data foundation, metadata contract or orchestration design must be refined.

Phase 5 should make explicit where AI is used and why. AI may support intent classification, metadata ranking, clarification, SQL generation, answer explanation, follow-up handling and safe-failure decisions. However, the MVP should not rely on the model alone for access control, metric logic, approved joins, validation, policy decisions or production assurance.

The MVP should include a lightweight user interface or interaction shell, but only to expose the flow, test user-facing behaviour and capture feedback. It should not become a full product-design exercise.

The phase should also introduce the minimum operational scaffolding needed to avoid rework later: environment separation, CI/CD discipline, prompt and model versioning, basic security controls, logging, tracing, cost and latency capture, error handling, feedback capture and a simple quality dashboard. These controls do not need to be production-grade, but they should make the MVP observable, testable and supportable.

Phase 5 should also state clearly what remains unresolved for production. Formal security assurance, full evaluation, load testing, release governance, incident management, monitoring thresholds, support model, data retention, privacy review, resilience, accessibility, adoption planning and production run-cost approval should be carried into later validation, pilot and production-readiness phases rather than implied by the MVP.

By the end of Phase 5, stakeholders should be able to decide whether the MVP provides enough evidence to proceed to formal validation, whether the scope should be constrained, whether the build or orchestration design needs remediation, whether the governed foundation needs refinement, or whether the initiative should pause.

The main output is a prototype evidence pack and a working MVP intended to progress into Phase 6 validation and then Phase 7 user testing. While Phase 6 may identify targeted improvements or corrective actions, the expectation is that the MVP built in Phase 5 will be used substantially as-is during pilot testing. Supporting outputs include the implemented flow, AI and model choices, validation approach, logs and traces, governance baseline, quality and cost findings, user feedback, known limitations, issue backlog and production-readiness gaps.

2. Phase overview

Phase 5 builds the bounded MVP from the approved Phase 4 orchestration design. It turns the designed question-to-answer flow into a working system that can be tested with realistic questions, controlled data, defined AI/model choices, validation checks, logging and user feedback.

The key discipline is to build an observable and testable MVP, not a polished demo. Phase 5 should make the approved flow real enough to test, but still bounded enough to remain safe, explainable and correctable.

2.1. Objective

The objective of Phase 5 is to build a bounded MVP and produce the evidence needed for formal validation, controlled user testing and production-readiness planning. It should confirm:

- **Implemented flow:** how the system moves from user question to metadata retrieval, AI/model routing, query generation, validation, execution, answering and feedback.
- **AI and model use:** where AI is used, which model patterns are selected, and where deterministic rules or human review are safer.
- **Data and metadata integration:** whether approved Phase 3 assets and Phase 4 metadata contracts can be consumed reliably at runtime.
- **Query and answer controls:** whether queries are bounded, validated and executed safely, validated and executed safely, and whether answers show sources, assumptions, caveats and safe-failure behaviour.
- **User interaction:** whether a lightweight interface supports realistic questions, follow-ups, clarification, refusals, answer review and feedback.
- **Operational scaffolding:** whether the MVP includes enough CI/CD, versioning, logging, tracing, basic security, cost capture, latency capture, error handling and quality monitoring to support validation and pilot use.
- **Production-readiness gaps:** what remains to be strengthened during validation, pilot and production-readiness phases.

2.2. Scope of the phase

Phase 5 should remain bounded to the approved users, questions, data assets, metadata sources, tools and orchestration design. It should implement enough of the system to test the intended flow without becoming a broad product build or uncontrolled expansion exercise.

In scope are the MVP application flow, AI/model calls, retrieval components, query generation, validation checks, execution hand-off, result interpretation, answer generation, lightweight UI, feedback capture, logging, tracing, cost and latency capture, basic security controls, deployment discipline and issue tracking.

The test for inclusion is simple: if an item is required to test the MVP safely, diagnose failures, support Phase 6 validation and remediation or prepare for Phase 7 user testing, it belongs in Phase 5. Issues exposed during build should be fixed, constrained, deferred or routed back rather than hidden in prompts or manual workarounds.

2.3. What this phase does not do

Phase 5 does not approve the system for production. A working MVP shows that the approved design can operate in a bounded way; it does not prove that the service is ready for broad release.

It also does not replace formal validation, redesign the governed foundation or deliver the full product experience. Issues with answer quality, security, access exposure, resilience, cost, scalability, Phase 3 artefacts or Phase 4 design assumptions should be tested, refined or routed into later phases rather than patched around during the build.

2.4. Expected duration and level of effort

The effort depends on delivery intent:

- For a POC, Phase 5 should prove that the main components can run together: approved data access, metadata retrieval, AI/model calls, query generation, validation, execution, answer generation and basic feedback.
- For an MVP, Phase 5 should produce a first usable version of the capability, with enough deployment discipline, versioning, validation, logging, cost/latency capture, feedback and quality monitoring to support Phase 6 validation and remediation and Phase 7 user testing. The phase is complete when the MVP can run the approved flow end-to-end, expose its limitations and provide enough evidence to proceed to formal validation.

2.5. Main participants and decision owners

Phase 5 is a build phase, but it should not be treated as an engineering-only activity. The MVP will only be useful if product, data, AI, security, evaluation and operating stakeholders help decide what is implemented, tested, constrained and carried forward.

Role	Main responsibility in Phase 5
Product owner	Confirms MVP scope, priorities, user-facing behaviour and trade-offs.
AI / solution architect	Owns the implemented orchestration flow, model/task choices, tool integration and design compromises.
AI / ML engineer	Builds model interactions, prompts, routing, retrieval, evaluation hooks and model/version controls.
Data / semantic owner	Confirms that metadata, metric definitions, caveats and approved artefacts are used correctly.
Analytics / data engineer	Implements data access, query execution, performance checks and data-layer integration.
Software engineer	Builds the application, APIs, deployment, UI and integration.
Security / governance lead	Reviews basic access controls, tool permissions, logging, data exposure and unresolved security gaps.
Evaluation owner	Defines technical tests, early answer-quality checks, issue taxonomy and evidence requirements for Phase 6.
Business SME	Tests questions, answers, caveats, clarifications / failure behaviour.
Operating owner	Reviews supportability, monitoring needs, incident implications, feedback handling and production-readiness gaps.

The key point is accountability. Phase 5 should not end with a working MVP that no one can explain, test, secure, support or improve.

3. Readiness decision and delivery implications

Phase 5 should not end with a statement that the MVP “works”. It should end with a clear decision on whether the implemented system provides enough evidence to move into formal validation.

The decision should be based on the working flow, model and tool behaviour, metadata grounding, query validation, answer quality, test results, logging, cost and latency observations, known limitations, governance baseline and unresolved production-readiness gaps.

A proceed decision does not mean the system is production-ready. It means the MVP is mature enough to be tested formally in Phase 6 and, if validated, used for controlled pilot testing in Phase 7.

For a POC, Phase 5 may be the final delivery stage if the objective is structured learning and feasibility assessment. If the output will be used by real users for controlled testing, decision support or pilot activity, it should proceed through formal validation before use.

3.1. Possible Phase 5 outcomes

Outcome	Meaning
Proceed to validation	The MVP is stable and evidenced enough to move into formal validation.
Proceed with constraints	Validation can start, but only for restricted users, questions, data, tools or answer types.
Remediate build	Implementation gaps must be fixed before validation.
Refine design	Phase 4 orchestration assumptions were incomplete, unsafe or impractical.
Refine foundation	Phase 3 assets or metadata are not usable enough at runtime.
Pause	Value, risk, cost, usability or ownership is not credible enough to continue.
Close POC	The POC has produced enough learning to stop, redesign or decide whether to invest in an MVP.

3.2. Minimum conditions to proceed

Phase 5 should proceed to formal validation only when the team can confirm:

- **The approved flow works end-to-end:** realistic questions can move through retrieval, AI/model routing, query generation, validation, execution, answer generation and feedback.
- **AI use is explicit:** model choices, deterministic controls and human-review points are documented.
- **Metadata grounding works:** the system can retrieve and apply the right metrics, dimensions, joins, filters, caveats and examples.
- **Queries are controlled:** unsafe, invalid, unauthorised or excessive queries are blocked or routed for review.
- **Answers are explainable:** responses show the relevant source, assumptions, caveats and safe-failure behaviour.

- **The MVP is testable and observable:** tests, logs, traces, validation outcomes, security/access events, cost, latency, errors and feedback are captured.
- **Governance is credible:** owners, change controls, issue routes, limitations and next actions are clear enough to support Phase 6 validation.

3.3. Common reasons to remediate, redesign or pause

The team should avoid moving into validation where any of the following apply:

- **The flow only works in demos:** realistic questions frequently fail, drift or require manual intervention.
- **Prompt patches hide design gaps:** prompts compensate for weak metadata, unclear logic or missing controls.
- **Retrieval is unreliable:** the system selects wrong, stale or incomplete metadata.
- **Validation is too shallow:** checks catch syntax errors but miss scope, access, joins, cost or plausibility.
- **Answers overstate certainty:** caveats, assumptions, limitations or refusals are missing or weak.
- **Logs are insufficient:** failures cannot be diagnosed, reproduced or evaluated.
- **Cost or latency is already fragile:** the MVP pattern is unlikely to support validation or pilot use.
- **Ownership is unclear:** no one owns defects, model behaviour, metadata fixes, quality issues or support.

3.4. How Phase 5 shapes later phases

Phase 5 provides the evidence base for formal validation, pilot planning and production-readiness work. Its outputs should not be treated as a demo record; they should become the test evidence, issue backlog and delivery constraints for later phases.

Later phase	What Phase 5 should provide
Phase 6 — Validation, assurance and remediation	Working MVP, test cases, logs, traces, validation results, quality findings, security gaps and failure examples.
Phase 7 — Controlled pilot	Usable MVP, supported scope, known limitations, feedback route, user-facing caveats and issue triage process.
Phase 8 — Production readiness and go-live	Production-readiness gaps, cost and latency evidence, support needs, governance gaps and hardening backlog.

A strong Phase 5 does not remove the need for validation. It makes validation possible.

4. MVP build activities overview

Phase 5 translates the approved design into an implemented MVP. The activities are ordered from scope confirmation and build setup through AI/model implementation, data and metadata integration, query control, answer generation, interface, testing and evidence packaging.

Governance runs across the sequence: decisions, changes, risks, model choices, validation gaps and production-readiness issues should be visible, owned and traceable throughout the build.

The activities should not be treated as a rigid waterfall. Build work will normally iterate as retrieval, prompts, validation, query execution, answer quality, cost, latency and logging issues are discovered. The important point is that changes should be visible, tested and governed rather than patched informally.

4.1. Activity sequence

Activity	Main question	Main output
1. Confirm MVP scope and build plan	What exactly will be built, tested and excluded?	Build scope, backlog and acceptance criteria.
2. Confirm MVP governance and decision rights	Who owns, validates and approves build decisions, changes, risks and evidence?	Role map, approval points, decision log and change route.
3. Set up environment, CI/CD and test harness	Can the MVP be built, deployed, tested and traced safely?	Environments, deployment path, test harness and trace setup.
4. Connect approved data, metadata and tools	Can the MVP access only approved assets and interfaces?	Working connections, permissions and access boundaries.
5. Implement metadata retrieval and grounding	Can the system retrieve and apply the right governed context?	Retrieval component, grounding tests and retrieval issues.
6. Implement AI/model routing and clarification	Which AI/model patterns are used, and when should the system clarify?	Routing, prompts, clarification rules and model choices.
7. Implement query generation, validation and execution	Can questions become bounded, validated and executable queries?	Query generation, validation gates and execution hand-off.
8. Implement result interpretation and answer generation	Can results become clear, caveated and explainable answers?	Answer generation, caveats and safe-failure behaviour.
9. Build lightweight UI and feedback capture	Can users test the flow without a full product UI?	Interaction shell, answer display and feedback capture.
10. Tune, test and package evidence	Are prompts, retrieval, model routing, validation and answers good enough for formal validation?	Test results, tuned behaviour, logs, traces, limitations, issue backlog and handover pack.

4.2. Activity logic

The build should follow the approved Phase 4 design, but not assume that the design is correct. Phase 5 is where retrieval quality, model choices, validation rules, metadata structure, query cost, answer format and failure behaviour are tested against implementation reality.

The MVP should therefore be built around observable steps. The team should be able to see which metadata was retrieved, which model or rule was used, what query was generated, which validations passed or failed, what was executed, how the answer was produced and what limitations remain.

4.3. Output need levels

Need level	Meaning
Mandatory	Required to complete Phase 5 and support a responsible validation decision.
Recommended	Strongly useful for MVP quality, but may be simplified for a narrow POC.
Optional / context-dependent	Needed only for higher-risk, regulated, complex or production-oriented paths.

For Phase 5, tuning and improvement should be treated as controlled build work: changes to prompts, retrieval, model routing, validation or answer behaviour should be tested, logged and owned.

4.4. Practitioner note

Phase 5 should not optimise for a smooth demo. It should optimise for a system that can be tested, challenged and improved. A brittle MVP with a good interface is less valuable than a narrower MVP that exposes its assumptions, logs its behaviour and makes the next delivery decision clear.

5. Core build activities

5.1. Confirm MVP scope and build plan

Purpose: Confirm what Phase 5 will build, test and exclude, and turn the approved Phase 4 design into a practical MVP build plan.

Activities

- Confirm the approved users, questions, answer types and exclusions.
- Confirm the Phase 3 assets, metadata sources and Phase 4 orchestration components required for the build.
- Define the MVP backlog across data, metadata, AI/model use, validation, UI, logging and testing.
- Agree the evidence required before the MVP can move to Phase 6 validation.
- Identify assumptions or risks that could force scope reduction during build.
- Confirm the MVP technology stack (application framework, orchestration layer, model gateway, retrieval store, query execution, logging/tracing, CI/CD and hosting environment).

Outputs and need level

Output	Description	Need level
MVP scope statement	Users, questions, data assets, tools, answer types and exclusions included in the build.	Mandatory
Build backlog	Prioritised implementation items across data, metadata, AI/model use, validation, UI, logging and testing.	Mandatory
Acceptance criteria	Minimum evidence required before the MVP can move to formal validation.	Mandatory
Known assumptions and risks	Design, data, metadata, access, cost, latency or usability risks that may affect build.	Mandatory
Prototype simplifications	Deliberate shortcuts allowed for the current stage, with limits and owners.	Recommended
Build timeline and dependencies	Delivery sequence, key dependencies and decision points.	Recommended

Watchpoints

Watchpoint	Delivery risk	Build response
Scope expands during build	The MVP becomes too broad to test or govern.	Keep to approved questions and record new ideas in backlog.
POC shortcuts enter MVP unchallenged	Weak controls become embedded in the product path.	Label simplifications clearly and decide whether they block validation.
Evidence needs are unclear	The MVP may work but still be hard to validate.	Define acceptance criteria before build starts.
Exclusions are vague	Stakeholders may assume unsupported questions are included.	Document exclusions and expected refusal or clarification behaviour.
Dependencies are unresolved	Build stalls or relies on manual workarounds.	Assign owners and route blockers back to the relevant phase.

Practitioner note

This activity should be straightforward if Phases 1 to 4 were completed properly, but it is still important. A short alignment meeting with product, data, AI, security, evaluation and operating stakeholders should be enough to confirm the MVP scope, technology choices, exclusions and evidence expectations. The build backlog should then be shaped jointly by the product owner and engineering lead so that delivery priorities remain aligned with validation needs.

5.2. Confirm MVP governance, ownership and decision rights

Purpose: Confirm how the ownership defined in earlier phases applies to the MVP build: who owns each component, who validates outputs, who approves changes and who decides whether defects or limitations block Phase 6 validation¹.

Activities

- Confirm named owners for the main MVP components: data access, metadata, retrieval, AI/model routing, prompts, validation, execution, UI, logging, testing and evidence.
- Confirm who validates key outputs, including generated queries, answer quality, caveats, refusals, logs, test results and production-readiness gaps.
- Agree approval points for material changes to scope, prompts, models, tools, validation rules, metadata, access assumptions, supported questions or answer formats.
- Define how defects, failed tests, risks, limitations and production-readiness gaps will be logged, prioritised, accepted or escalated.
- Confirm who owns the Phase 6 validation handover, including evidence quality, unresolved issues and known constraints.
- Identify which governance elements are MVP-ready and which must be strengthened before pilot or production.
- Create or confirm an MVP governance register covering ownership, validation responsibilities, decision rights, issue route, change log, risk acceptance and production-readiness gaps.

Outputs and need level

Output	Description	Need level
MVP governance register	Ownership, validation responsibilities, decision rights, issue route, change log, risk acceptance and governance gaps.	Mandatory for MVP / Recommended for POC
Component ownership map	Named owners for build components and evidence areas.	Mandatory
Approval route	How material changes to scope, prompts, models, tools, validation, metadata or access are approved.	Mandatory
Phase 6 handover owner	Named owner accountable for evidence quality and validation handover.	Mandatory

¹ For a POC, this activity may be lightweight. For an MVP intended to enter validation and user testing, governance should be documented clearly enough to support formal testing, controlled change, issue ownership and later production readiness

Production governance gaps	Governance items still required before pilot or production.	Recommended
-----------------------------------	---	-------------

Watchpoints

Watchpoint	Delivery risk	Build response
Governance is treated as later-phase work	Build decisions become informal and hard to validate.	Define minimum decision rights before implementation starts.
Prompt or model changes are undocumented	Test results become difficult to reproduce or trust.	Log prompt, model and routing changes with owner and reason.
No one owns validation failures	Issues remain unresolved or are hidden in workarounds.	Assign each failure to a named owner and decision route.
Business validation is missing	Answers may look technically correct but be wrong in business terms.	Involve business and semantic owners in answer review.
Security assumptions are accepted silently	Access or exposure issues may become embedded in the MVP.	Record assumptions, unresolved risks and approval status.

Practitioner note

This activity is often overlooked because it can feel administrative during build, but it is critical for later validation, pilot and production readiness. Governance creates clear ownership, and clear ownership keeps senior stakeholders engaged when trade-offs, defects, risks and limitations need decisions rather than informal workarounds.

5.3. Set up environment, CI/CD and test harness

Purpose: Create the technical scaffolding needed to build, test, trace and change the system safely. For a POC this may be minimal; for an MVP it should be strong enough to support validation, user testing and later handover.

Activities

- Set up development and test environments with approved access to data, metadata, models and tools.
- Configure version control, deployment route and CI/CD discipline appropriate to the delivery stage.
- Set up secrets, credentials and access controls for model, data and tool connections.
- Create a repeatable test harness for success, failure, unsafe and regression cases.
- Configure logging, tracing, cost and latency capture from the start.

Outputs and need level

Output	Description	Need level
MVP environments	Controlled development and test environments with approved access.	Mandatory
Versioning and deployment route	Controlled route for code, prompts, configuration and deployment.	Mandatory
Secrets and access setup	Controlled handling of credentials, keys and tool permissions.	Mandatory

Test harness	Repeatable tests for success, failure, unsafe and regression cases.	Mandatory
Observability baseline	Logs, traces, cost and latency signals needed for validation.	Mandatory for MVP / Recommended for POC

Watchpoints

Watchpoint	Delivery risk	Build response
Local prototype only	The MVP cannot be reproduced, tested or handed over.	Use a controlled environment and documented setup.
No prompt or config versioning	Behaviour changes cannot be explained or reproduced.	Version prompts, model settings and configuration.
Weak test harness	The MVP works in demos but fails under realistic questions.	Include success, failure, unsafe and regression cases.
Logging added too late	Failures cannot be diagnosed or used for validation.	Capture traces from the start of implementation.
Secrets handled informally	Keys, credentials or permissions may be exposed.	Use controlled secrets and access management.
Cost and latency ignored	The chosen design may be unusable at validation or pilot scale.	Capture basic cost and response-time signals early.

Practitioner note

Note 1: This is often where shortcuts are taken because the team wants to get to the visible prototype quickly. That can be acceptable for a narrow POC. For an MVP, the pressure is often to demonstrate value quickly while also educating stakeholders on why delivery discipline matters. Finding the right balance can be difficult, but the basics should still be in place early. Retrofitting version control, deployment discipline, logging, tracing, secrets management, test harnesses and issue tracking later is usually slower and riskier than setting them up properly at the start.

Note 2: Developers should also be able to create a small local or isolated development environment through a documented script, configuration or infrastructure template. This allows them to work independently, reproduce issues and test changes without relying on a shared fragile setup.

5.4. Connect approved data, metadata and tools

Purpose: Connect the MVP to the approved data assets, metadata sources and tools required to run the Phase 4 flow. The aim is to prove controlled runtime access, not to broaden the data estate.

Activities

- Connect the MVP to approved Phase 3 data assets, using scoped read-only access where possible.
- Connect to the metadata sources needed for runtime grounding, including metric, dimension, join, caveat and example artefacts.
- Connect the required tools, including metadata retrieval, query validation, query execution, logging, tracing and feedback capture.

- Confirm that access boundaries, permissions, row/column controls, masking assumptions and blocked assets are enforced or clearly documented.
- Test the connected data, metadata and tools through representative approved questions to confirm they support the intended end-to-end flow

Outputs and need level

Output	Description	Need level
Approved connection list	Data assets, metadata sources and tools connected to the MVP.	Mandatory
Access boundary record	What the MVP can and cannot access, including blocked sources, fields or actions.	Mandatory
Metadata availability check	Confirmation that required runtime metadata is available and usable.	Mandatory
Connection test results	Evidence that data, metadata and tool connections work for supported questions.	Mandatory
Unresolved access or tooling gaps	Issues that may block validation, pilot or production readiness.	Mandatory

Watchpoints

Watchpoint	Delivery risk	Build response
Unapproved data is connected	The MVP tests a flow that cannot be governed later.	Connect only approved assets and log exclusions.
Metadata is available but not usable	Retrieval or answer generation fails at runtime.	Test metadata against supported questions.
Tool permissions are too broad	The MVP can query, expose or execute beyond scope.	Restrict tools and document permissions.
Access controls are assumed	Security gaps remain hidden until validation.	Test role, row, column and masking assumptions early.
Connection tests are too technical	The system connects but cannot answer real questions.	Test connections through the full question-to-answer flow.

Practitioner note

Access metadata may define the user's security level, role, region entitlement or permitted dimensions, but it should drive an enforced policy before execution. The MVP should test that generated queries, semantic-layer requests or API calls are constrained by those entitlements before restricted results reach the model or answer layer. Metadata should inform access; policy should enforce it.

5.5. Implement metadata retrieval and grounding

Purpose: Build the retrieval path that provides the model with the governed context it needs to answer safely. The aim is to test whether approved metrics, dimensions, joins, filters, caveats, examples and access rules can be found and applied at runtime.

Activities

- Implement retrieval over the approved metadata sources, including metric, dimension, join, caveat, example and access-control artefacts.

- Define how retrieved context is ranked, filtered, versioned and passed into prompts or tools.
- Test retrieval against supported MVP questions, ambiguous questions and out-of-scope questions.
- Confirm that retrieved context is limited to approved, current and user-permitted artefacts.
- Record retrieval failures, weak matches, missing metadata and grounding issues for remediation or backlog.

Outputs and need level

Output	Description	Need level
Retrieval implementation	Working retrieval path over approved metadata sources.	Mandatory
Grounding rules	Rules for what metadata is retrieved, ranked, filtered and passed to the model.	Mandatory
Retrieval test results	Evidence that supported questions retrieve the right context.	Mandatory
Metadata gap log	Missing, weak, stale or unusable metadata discovered during testing.	Mandatory
Retrieval quality measures	Early view of retrieval accuracy, weak matches, false matches and failure cases.	Recommended

Watchpoints

Watchpoint	Delivery risk	Build response
Retrieval finds the wrong context	The model may generate valid but incorrect answers.	Test retrieval against known questions and expected metadata.
Too much context is retrieved	Prompts become noisy, costly or inconsistent.	Limit context by scope, ranking, version and relevance.
Stale or draft metadata is used	The MVP may apply outdated definitions or caveats.	Retrieve only approved/current artefacts unless explicitly testing drafts.
Access rules are not applied to metadata	Users may see or infer restricted sources, fields or definitions.	Filter metadata by user entitlement and approved scope.
Missing metadata is hidden by prompts	The model compensates instead of failing safely.	Log gaps and clarify, refuse or route for remediation.

Practitioner note

This activity should be straightforward if Phase 3 produced usable metadata artefacts and Phase 4 defined the retrieval pattern. The main task is to implement and test the retrieval path, not redesign the semantic layer. However, weak retrieval can quietly break the MVP by grounding the model in the wrong metric, source, join, filter or caveat.

5.6. Implement AI/model routing and clarification

Purpose: Implement how the MVP decides which AI/model pattern, rule or tool path should handle each request, and when the system should clarify instead of proceeding.

Activities

- Implement routing rules for supported domains, question types, model tasks and tool paths.
- Configure model choices for intent handling, metadata ranking, clarification, query generation, answer generation and safe-failure support.
- Define when deterministic rules should override model judgement, especially for scope, access, validation and refusal decisions.
- Implement clarification behaviour for ambiguous metrics, missing filters, unclear time periods, unsupported dimensions or unsafe requests.
- Log routing decisions, model choices, clarification triggers, fallback paths and unresolved routing issues.

Outputs and need level

Output	Description	Need level
Routing implementation	Working routing logic for supported domains, questions, models and tools.	Mandatory
Model/task map	Which model, rule or deterministic control is used for each task.	Mandatory
Clarification rules	Conditions that trigger clarification before query generation or answer generation.	Mandatory
Fallback and refusal path	What happens when routing fails, context is missing or the request is unsupported.	Mandatory
Routing and clarification test results	Evidence that routing and clarification work for supported, ambiguous and unsupported questions.	Mandatory

Watchpoints

Watchpoint	Delivery risk	Build response
One model handles everything	Cost, latency and failure risk increase unnecessarily.	Route tasks by complexity, risk and evidence.
Routing is hidden in prompts	Behaviour becomes hard to test, reproduce or govern.	Make routing rules explicit and logged.
Clarification is too weak	The system guesses when it should ask.	Define mandatory clarification triggers.
Clarification is too frequent	Users are blocked by unnecessary questions.	Test against realistic MVP questions and refine thresholds.
Model choice is not evidenced	The team cannot justify cost, quality or safety trade-offs.	Capture test results, cost and latency by model/task.

Practitioner note

This activity should not become model shopping. In many MVPs, metadata retrieval and ranking can use embeddings, rules or smaller models; SQL generation often needs a medium model, with stronger models reserved for complex queries; answer generation can often use a medium

model if the result and caveats are well grounded. Stronger models should be justified by test results, not selected by default. Routing, access, validation and refusal should remain controlled by explicit rules and governed logic rather than unconstrained model judgement.

5.7. Implement query generation, validation and execution

Purpose: Implement the path that turns a grounded user question into a bounded, validated and executable query. The aim is not only to generate SQL, but to prove that generated queries stay within approved scope, access rules, cost limits and validation controls.

Activities

- Implement query generation using approved metadata, examples, metric logic, dimensions, joins, filters and access context.
- Apply rule-based validation before execution, including approved assets, joins, filters, permissions, row limits, blocked actions and cost thresholds.
- Execute only validated read-only queries through the approved warehouse, semantic layer or API route; editing, write-back and destructive actions should be blocked.
- Apply query limits and guardrails, such as row limits, timeout limits, scan limits, aggregation rules and restricted joins.
- Capture generated queries, validation outcomes, execution status, server-side performance, errors, latency, cost signals and result shape.
- Test supported questions, unsafe questions, ambiguous questions, access-restricted questions and expected validation failures.

Outputs and need level

Output	Description	Need level
Query generation implementation	Working generation path using approved metadata and examples.	Mandatory
Rule-based validation gates	Deterministic checks applied before query execution.	Mandatory
Controlled execution path	Approved read-only route with limits, blocked actions and access controls.	Mandatory
Query performance evidence	Server-side runtime, timeout, scan, row-count and cost observations.	Mandatory for MVP / Recommended for POC
Query and validation test results	Evidence from supported, unsafe, restricted and failed-query cases.	Mandatory
Query issue log	Defects, invalid SQL, wrong joins, excessive cost, access failures and fixes.	Mandatory

Watchpoints

Watchpoint	Delivery risk	Build response
Valid SQL is treated as correct SQL	The query runs but uses the wrong metric, grain, join or filter.	Validate business logic, not only syntax.

Validation is model-led only	Unsafe queries may be approved by the same system that generated them.	Use rule-based checks before execution.
Access filters are added by the model	User entitlements may be applied inconsistently.	Use governed access metadata and enforced policy logic.
Query execution is too permissive	The MVP can scan, join, edit or expose more than intended.	Restrict to read-only approved assets, limits and blocked actions.
Server performance is ignored	The query may work once but fail under validation or pilot conditions.	Track server runtime, scan volume, row counts, timeouts and cost.

Practitioner note

This activity is where many T2D prototypes look impressive but become unsafe. The model may generate syntactically valid SQL, but that does not mean the query is correct, authorised, affordable or safe to run. For an MVP, query execution should be read-only, rule-validated, bounded by limits and measured for server-side performance before it is trusted for validation or pilot use.

5.8. Implement result interpretation and answer generation

Purpose: Implement how query results are turned into clear, grounded and caveated answers. The aim is to make the answer useful to the user without losing the source, calculation, assumptions, limitations or safe-failure behaviour.

Activities

- Interpret query results using the executed query, result set, metric definitions, filters, caveats and access context.
- Generate answers in the agreed MVP format, including source, period, metric, filters, assumptions and limitations where relevant.
- Apply answer rules for caveats, uncertainty, empty results, partial results, blocked results and failed queries.
- Test answer outputs against evidence rules, including checks for unsupported explanations, causal claims, recommendations, confidence wording and missing caveats
- Capture answer examples, reviewer comments, answer-quality issues and required prompt/model improvements.

Outputs and need level

Output	Description	Need level
Answer generation implementation	Working answer path using query results, metadata and caveats.	Mandatory
Answer format rules	Required structure for numbers, tables, narrative answers, caveats, refusals and clarifications.	Mandatory
Caveat and limitation handling	Rules for when caveats, assumptions, uncertainty or safe-failure wording must be shown.	Mandatory

Answer-quality examples	Examples of good, partial, failed, caveated and refused answers.	Mandatory
Answer issue log	Incorrect, overstated, unclear, missing-caveat or unsupported-answer issues.	Mandatory

Watchpoints

Watchpoint	Delivery risk	Build response
Answers sound better than the evidence	Users may over-trust partial or weak results.	Require caveats, assumptions and source visibility.
The model adds unsupported interpretation	The answer may imply causes, advice or certainty not in the data.	Restrict answer generation to result, metadata and approved context.
Caveats are dropped	Known limitations disappear from the user-facing answer.	Treat mandatory caveats as required answer content.
Empty or partial results are mishandled	The system may present missing data as a real zero or complete answer.	Define specific empty, partial and failed-query behaviours.
Confidence wording is unclear	Users may misread technical uncertainty as business certainty.	Use confidence wording only where defined and understandable.

Practitioner note

Answer generation should be tested with questions that tempt the model to overreach, such as “why did revenue fall?”, “what should we do next?” or “are we performing well?”. Unless the retrieved evidence includes approved causal logic, benchmarks or recommendations, the MVP should return the result, caveat the limitation and explain that the available data does not support a cause, judgement or recommendation. This behaviour should be tested and logged, not left to prompt wording alone.

5.9. Build lightweight UI and feedback capture

Purpose: Build the minimum user-facing interface needed to test the MVP flow, answer presentation and feedback route. The aim is not to design the final product experience, but to make the system usable enough for validation and controlled user testing.

Activities

- Build a lightweight interaction shell for asking questions, receiving answers, seeing caveats and reviewing safe-failure responses.
- Display the answer elements needed for trust, including result, source, metric, filters, assumptions, caveats and limitations where relevant.
- Support MVP interaction behaviours, including clarification, refusal, follow-up questions and visible error handling.
- Capture structured feedback, corrections, defects and questions that could not be answered.
- Ensure feedback links back to traces, generated queries, retrieved metadata, model versions and issue tracking.

Outputs and need level

Output	Description	Need level
Lightweight UI	Simple interface or interaction shell for MVP testing.	Mandatory
Answer display pattern	How results, sources, filters, caveats, limitations and safe-failure messages are shown.	Mandatory
Feedback capture route	How testers record usefulness, trust, corrections, defects and missing answers.	Mandatory
Trace-to-feedback link	Ability to connect feedback to trace ID, query, metadata, model version or issue.	Mandatory for MVP / Recommended for POC
UI issue log	Usability, explanation, trust, display or feedback-capture issues.	Recommended

Watchpoints

Watchpoint	Delivery risk	Build response
UI becomes the project	Effort shifts from testing the T2D flow to building product polish.	Keep the interface simple and focused on evidence capture.
Answers hide caveats	Users may over-trust results.	Make caveats, filters and limitations visible.
Feedback is unstructured	Issues cannot be linked to root cause or validation evidence.	Link feedback to traces, queries, metadata and model versions.
Follow-ups are unclear	Users may not know whether context was inherited or reset.	Make key inherited filters or context visible where useful.
Errors are too technical	Users cannot tell whether to retry, clarify or escalate.	Use clear error, refusal and clarification messages.

Practitioner note

The Phase 5 UI should be designed with the target users, but it should remain deliberately lightweight. The interface should fit the usage context: a web UI or embedded app may suit desk-based users, while mobile or field users may need a different interaction pattern, including voice where appropriate. Security should remain central: the UI must not expose restricted data, hidden prompts, raw traces or query details to users who are not entitled to see them.

5.10. Tune, test and package evidence

Purpose: Improve the MVP until it is good enough to enter formal validation, and package the evidence needed for Phase 6. The aim is controlled tuning, not informal prompt patching.

Activities

- Run the MVP against supported questions, ambiguous questions, unsafe questions, access-restricted questions and expected failure cases.
- Tune versioned prompts, retrieval, model routing, validation rules and answer behaviour based on test evidence.
- Retest changed behaviour and record before/after results, prompt, model, retrieval and validation versions, and remaining defects.
- Prioritise issues as blocking, must-fix, accepted limitation, deferred improvement or production-readiness gap.

- Package the evidence needed for Phase 6, including test results, logs, traces, examples, known limitations, issue backlog and governance decisions.

Outputs and need level

Output	Description	Need level
MVP test results	Evidence from supported, ambiguous, unsafe, restricted and failure cases.	Mandatory
Tuning record	Changes to prompts, retrieval, model routing, validation or answer behaviour, with rationale and evidence.	Mandatory
Retest evidence	Before/after results showing whether changes improved behaviour or introduced regressions.	Mandatory
Issue backlog	Defects, limitations, risks, accepted gaps and production-readiness items.	Mandatory
Phase 6 evidence pack	Working MVP evidence required for formal validation planning.	Mandatory

Watchpoints

Watchpoint	Delivery risk	Build response
Tuning becomes prompt hacking	Behaviour improves in demos but becomes hard to reproduce or govern.	Log changes, versions, rationale and test evidence.
Tests only cover success cases	The MVP looks ready but fails on ambiguity, access or unsafe requests.	Include failure, refusal, clarification and restricted-access cases.
Fixes are not retested	Improvements may break previously working questions.	Run regression checks after material changes.
Issues are hidden to proceed	Phase 6 starts with unknown or unmanaged risks.	Classify issues and record owners, severity and next action.
Evidence is incomplete	Validation teams cannot reproduce or trust the MVP behaviour.	Package traces, tests, examples and decisions before handover.

Practitioner note

The final Phase 5 activity should make the MVP better, not just document what happened. Tuning is valid when it is controlled, tested and recorded. It becomes dangerous when prompt changes, model switches or validation relaxations are made informally to make a demo pass. Phase 6 should receive an MVP whose behaviour can be reproduced, challenged and validated.

6. Consolidated Phase 5 outputs

Phase 5 should end with a working MVP and a clear evidence pack. The output should not be a demo narrative; it should show what was built, how it behaves, what was tested, what changed during tuning, and what remains unresolved before validation, pilot or production.

6.1. MVP evidence pack

Output	Purpose
Working MVP	Demonstrates that the approved Phase 4 flow can run end-to-end in a bounded way.
Implemented flow record	Shows how user questions move through retrieval, model routing, query generation, validation, execution and answer generation.
AI/model decision record	Documents where AI is used, which model patterns were selected and where deterministic controls apply.
Test results	Shows supported, ambiguous, unsafe, access-restricted and failure cases.
Logs and traces	Allows failures, answers, generated queries, retrieved metadata and model calls to be diagnosed.
Quality and cost findings	Summarises early answer-quality, latency, server performance and cost observations.
Known limitations and issue backlog	Captures defects, accepted limitations, deferred items and production-readiness gaps.
Governance record	Confirms owners, decision rights, approval route, change log and Phase 6 handover ownership.

6.2. Output quality test

The Phase 5 outputs should be tested before handover. A useful evidence pack should be:

Quality test	Question
Complete	Does it show what was built, tested, changed and excluded?
Reproducible	Can Phase 6 replay or inspect the same prompts, models, metadata, queries, tests and versions?
Diagnostic	Can failures be traced to retrieval, model routing, SQL generation, validation, execution or answer generation?
Governed	Are owners, decisions, accepted risks and change records clear?
Bounded	Are scope, access, tool limits, simplifications and known gaps explicit?
Actionable	Does the backlog clearly separate blockers, validation issues, pilot risks and production-readiness gaps?

If these tests fail, Phase 5 may have produced a working MVP, but not enough evidence to support formal validation.

7. Exit criteria and handover

Phase 5 should close only when the MVP is usable, observable and evidenced enough to enter formal validation. The handover should make clear what was built, what was tested, what remains weak and what must be validated before controlled user testing.

7.1. Required exit outputs

Before Phase 5 closes, the team should have:

Output	Purpose
Working MVP	Bounded implementation of the approved question-to-answer flow.
MVP evidence pack	Test results, traces, query examples, answer examples, tuning records and limitations.
Governance record	Owners, decision rights, change log, risk acceptance and Phase 6 handover owner.
Test and tuning results	Evidence of supported, ambiguous, unsafe, access-restricted and failure cases.
Issue backlog	Defects, blockers, accepted limitations, deferred items and production-readiness gaps.
Cost and latency findings	Early view of model cost, query performance, response time and operational constraints.
Validation recommendation	Decision on whether to proceed to Phase 6, proceed with constraints, remediate, redesign, refine foundation, close POC or pause.

7.2. Handover to later phases

Later phase	What Phase 5 should hand over
Phase 6 — Validation, assurance and remediation	Working MVP, evidence pack, test cases, traces, prompt/model versions, generated queries, validation outcomes, answer examples, known limitations and governance decisions.
Phase 7 — Controlled pilot	Validated MVP scope, user-facing limitations, feedback route, issue triage process, support assumptions and pilot constraints.
Phase 8 — Production readiness and go-live	Production-readiness gaps, hardening backlog, cost and latency evidence, governance gaps, support needs and operating assumptions.

A POC may close at Phase 5 if its purpose was only structured learning. An MVP intended for user testing, decision support or pilot use should not bypass Phase 6 validation.

7.3. Exit decision wording

Suggested wording for the Phase 5 exit decision:

Phase 5 is complete. The MVP has / has not produced enough evidence to proceed to Phase 6 validation. The agreed decision is [proceed / proceed with constraints / remediate build / refine design / refine foundation / close POC / pause]. The main constraints, accepted limitations and required next actions are [summary].

7.4. Practitioner note

A working MVP should not be confused with a validated product. Phase 5 proves that the system can be built and tested in a bounded way; Phase 6 proves whether it is accurate, safe and controlled enough to place in front of users.